

Big Data Paper

By Alexander Richin

Date: March 7, 2017

Bigtable:

A Distributed Storage System for Structured Data

Andrew Pavlov

&

Brown University

Pavlo@cs.brown.edu

A Comparison of Approaches to Large-Scale Data Analysis

Fay Chang

&

Google, Inc.

Fay@google.com

One Size Fits All - An Idea Whose Time Has Come and Gone (2005)

Michael Stonebraker

Brown University

Database Group

Main Idea – Bigtable

- A compressed, high performance, & proprietary data storage system
- Built on:
 - Google File System
 - Chubby Lock Service
 - SSTable (log-structured storage like LevelDB)
 - Etc.
- May 6, 2015
 - Public version of **Bigtable** available as service

Bigtable Implementation

- 3 Major Components - Bigtable Implementation
 - a. A library that is linked to every client
 - b. One master server
 - c. Many tablet servers
- Tablet Location
 - a. Three-level hierarchy analogous to store info.
- Tablet Assignment
 - a. Each tablet assigned to 1 tablet server at a time
- Tablet Serving
 - a. Persistent state of tablet stored in GFS
- Compactions
 - a. Shrinks memory usage of tablet server
 - b. Reduces amount of data read from commit log if server dies

Analysis of Bigtable

- Description - A distributed system for storing structured data @ Google
- What users like:
 - Performance & high availability
 - Scale capacity of clusters via adding more machines as resource demands change
- Difficult for new users to get adjusted to
 - If you are used to Relational Databases
- However, Google has demonstrated Bigtable's success
- New Features (coming soon!):
 - Support secondary indices & infrastructure for building cross-data-center replicas
 - Multiple master replicas
 - Serviced to product groups
 - No need for maintenance of clusters
- As service clusters scale:
 - Deal w/ more resource-sharing issues
 - API, Implementation

Main Idea (MapReduce vs Parallel DBMS)

- “Cluster Computing” - Constructing a data structure via large # of low end processors rather than smaller sets of high-end servers.
- MapReduce:
 - Attractive, simple, expresses sophisticated distributive programs
 - Data can be in arbitrary format
- Parallel DBMS
 - Vertica, Oracle, Teradata
 - Robust high-computing platforms
 - Data conforms to well-defined schema
- Main Idea:
 - “The purpose... to consider these choices, and the trade-offs they entail”

Implementation

● Architectural Elements

- Schema Support
- Parallel DBMS req. data to fit in relational paradigm
 - Rows / Columns
- MR does not
 - Most popular implementation of MR framework - Hadoop
- Indexing
 - MR does not have built in indexing
 - Parallel DBMS - hash / B-tree
- Data Distribution
 - Parallel - use data distribution & location to advantage
 - MR - must do manually
- Node configuration
 - Hadoop, Vertica, DBMS-X
 - Implemented on 100 node cluster

Analysis

● MapReduce

- Very simple to use
- Better for smaller projects
- Most flexible
- Better Fault Tolerance
- Well suited for development env. with:
 - Small # of programmers
 - Limited application domain
- Not appropriate for long-term / larger-sized projects
 - Lack of constraints

● Parallel DBMS

- More complex
- Harder to understand
- Not as flexible as MR
- Better Data Distribution

Bigtable vs. MapReduce / Parallel DBMS

- Comparisons:
 - All three deal with having a structured data storage
 - Deal with large amounts of data
 - Each is good in its own way
 - Depends upon what you use it for
 - Compress files and store them in an efficient manner
 - Can perform similar tasks to each other
 - Each is specified for its own purpose and has its trade offs
 - Some perform certain tasks more efficiently than others
 - One is not better than the other
 - Implementation is similar

Michael Stonebraker's Main Idea

-
- “Dead on our feet” (80’s & 90’s)
 - B/c we believed:
 - “One Size Fits All”
 - Engines are too diverse
 - We have moved on
 - “One size fits none”
 - We live in very interesting times
 - Good for DBMS researchers
 - New ideas (NVRAM, LLVM, Xeon/Phi)
 - People should explore DBMS

Advantages vs. Disadvantages

Bigtable Advantages:

- Substantial flexibility
 - Designing own data module
- Removes:
 - Bottlenecks
 - Inefficiencies

Parallel DBMS

- Can handle massive amounts of data
- Not so flexible
 - Follows structure
- Difficult to understand
 - If you're new

Large-Scale Data Analysis:

- Uses a significant amount of less energy
- Large user community
- Easy to set up

MapReduce:

- Simplest to use
- Most flexible
- Very efficient
- Only good for smaller projects
 - Less programmers